

Graphs

May 22, 2018

1 CIS 545 Homework 2: Graphs

```
In [1]: # Execute this once, the first time you run
!pip install networkx

# Disable Python warning messages - you should probably only run this before submission

import warnings
warnings.filterwarnings('ignore')
```

Requirement already satisfied: networkx in /opt/conda/lib/python3.6/site-packages

Requirement already satisfied: decorator>=4.1.0 in /opt/conda/lib/python3.6/site-packages (from

1.1 Step 2.1 Spark Setup

```
In [2]: # TODO: Connect to Spark as per Step 2.1
from pyspark.sql import SparkSession
from pyspark.sql.types import *
import pyspark.sql.functions as F

spark = SparkSession.builder.appName('Graphs-HW2').getOrCreate()

In [3]: # Load some dummy data, which should be overwritten in Step 2.2

answers_sdf = spark.createDataFrame([{'from_node': 123, 'to_node': 456},\
                                     {'from_node': 456, 'to_node': 789},
                                     {'from_node': 456, 'to_node': 890}])
comments_answers_sdf = spark.createDataFrame([{'from_node': 123, 'to_node': 456}])
comments_questions_sdf = spark.createDataFrame([{'from_node': 123, 'to_node': 456}])

graph_sdf = spark.createDataFrame([{'from_node': 123, 'to_node': 456}])
```

1.2 Step 2.2 Loading

```
In [4]: # TODO: load data as per Step 2.2
answers_sdf = spark.read.load('sx-stackoverflow-a2q.txt', format="text")
```

```

comments_answers_sdf = spark.read.load('sx-stackoverflow-c2a.txt', format="text")
comments_questions_sdf = spark.read.load('sx-stackoverflow-c2q.txt', format="text")

```

```

In [5]: # You may add as many cells as you like here.
        # Use Insert / Insert Cell Below

```

1.3 Step 2.2 Results

```
In [6]: answers_sdf.count()
```

```
Out[6]: 17823525
```

```
In [7]: answers_sdf.show()
```

```

+-----+
|          value|
+-----+
|  9 8 1217567877|
|  1 1 1217573801|
| 13 1 1217606247|
| 17 1 1217617639|
| 48 2 1217618182|
| 17 1 1217618239|
| 19 9 1217618357|
|13 23 1217618560|
|13 11 1217618799|
|23 23 1217619360|
|35 33 1217620542|
|39 33 1217620597|
|43 22 1217620971|
|17 32 1217621272|
|39 40 1217621416|
|37 40 1217621670|
|45 45 1217621917|
|17 17 1217622124|
|49 13 1217623079|
|13 23 1217623216|
+-----+
only showing top 20 rows

```

```
In [8]: answers_sdf.printSchema()
```

```

root
|-- value: string (nullable = true)

```

```
In [9]: comments_answers_sdf.count()
```

```
Out[9]: 25405374
```

```
In [10]: comments_answers_sdf.show(10)
```

```
+-----+
|          value|
+-----+
|    1  91 1220713630|
|    3  91 1220713792|
|   380 350 1220715736|
|4642 2257 1220734307|
|4642 1324220 1220...|
|2495 4285 1220736640|
|4642 4893 1220737355|
|2515 4903 1220738560|
|2515 4893 1220739071|
|   199 199 1220741079|
+-----+
only showing top 10 rows
```

```
In [11]: comments_answers_sdf.printSchema()
```

```
root
|-- value: string (nullable = true)
```

```
In [12]: comments_questions_sdf.count()
```

```
Out[12]: 20268151
```

```
In [13]: comments_questions_sdf.show(10)
```

```
+-----+
|          value|
+-----+
|4550 4550 1220729190|
|   242 184 1220733503|
|4213 4946 1220768149|
|    91  91 1220768295|
|2658 1874 1220771891|
|4035 1874 1220773037|
|2257 4489 1220802041|
|   577 577 1220834891|
|4489 4489 1220853536|
|   828 2783 1220854143|
+-----+
```

only showing top 10 rows

```
In [14]: comments_questions_sdf.printSchema()
```

```
root
|-- value: string (nullable = true)
```

1.4 Step 2.3

```
In [15]: # TODO: wrangling work in Step 2.3. Add as many Cells as you need
#answers_sdf.createOrReplaceTempView('answers_view')
#answers_sdf=spark.sql('select split(value," ")[0] As from_node, split(value," ")[1] As to_node, split(value," ")[2] As edge_type from answers_sdf')
answers_sdf=answers_sdf.withColumn('new_col',F.lit('answer'))
answers_sdf1=answers_sdf.select(F.split(answers_sdf.value,' ')[0].alias('from_node').cast('string'),F.split(answers_sdf.value,' ')[1].alias('to_node').cast('string'),F.split(answers_sdf.value,' ')[2].alias('edge_type').cast('string'))
answers_sdf1=answers_sdf1.select(answers_sdf1.from_node,answers_sdf1.to_node,answers_sdf1.edge_type)
#answers_sdf1.show(6)
answers_sdf=answers_sdf1
## Step 2.3 Results
```

```
In [16]: #answers_sdf.createOrReplaceTempView('answers_view')
comments_answers_sdf=comments_answers_sdf.withColumn('new_col',F.lit('comment-on-answer'))
comments_answers_sdf1=comments_answers_sdf.select(F.split(comments_answers_sdf.value,' ')[0].alias('from_node').cast('string'),F.split(comments_answers_sdf.value,' ')[1].alias('to_node').cast('string'),F.split(comments_answers_sdf.value,' ')[2].alias('edge_type').cast('string'))
comments_answers_sdf1=comments_answers_sdf1.select(comments_answers_sdf1.from_node,comments_answers_sdf1.to_node,comments_answers_sdf1.edge_type)
#comments_answers_sdf1.show(6)
comments_answers_sdf=comments_answers_sdf1
```

```
In [17]: #answers_sdf.createOrReplaceTempView('answers_view')
comments_questions_sdf=comments_questions_sdf.withColumn('new_col',F.lit('comment-on-question'))
comments_questions_sdf1=comments_questions_sdf.select(F.split(comments_questions_sdf.value,' ')[0].alias('from_node').cast('string'),F.split(comments_questions_sdf.value,' ')[1].alias('to_node').cast('string'),F.split(comments_questions_sdf.value,' ')[2].alias('edge_type').cast('string'))
comments_questions_sdf1=comments_questions_sdf1.select(comments_questions_sdf1.from_node,comments_questions_sdf1.to_node,comments_questions_sdf1.edge_type)
#comments_questions_sdf1.show(6)
comments_questions_sdf=comments_questions_sdf1
```

```
In [18]: graph_sdf=answers_sdf.unionAll(comments_questions_sdf).unionAll(comments_answers_sdf)
```

```
In [19]: answers_sdf.count()
```

```
Out[19]: 17823525
```

```
In [20]: answers_sdf.show(5)
```

```
+-----+-----+-----+
|from_node|to_node|edge_type|
+-----+-----+-----+
|          9|          8|    answer|
```

1	1	answer
13	1	answer
17	1	answer
48	2	answer

+-----+-----+-----+

only showing top 5 rows

```
In [21]: answers_sdf.printSchema()
```

```
root
 |-- from_node: integer (nullable = true)
 |-- to_node: integer (nullable = true)
 |-- edge_type: string (nullable = false)
```

```
In [22]: comments_answers_sdf.count()
```

```
Out[22]: 25405374
```

```
In [23]: comments_answers_sdf.show(5)
```

1	91	comment-on-answer
3	91	comment-on-answer
380	350	comment-on-answer
4642	2257	comment-on-answer
4642	1324220	comment-on-answer

+-----+-----+-----+

only showing top 5 rows

```
In [24]: comments_answers_sdf.printSchema()
```

```
root
 |-- from_node: integer (nullable = true)
 |-- to_node: integer (nullable = true)
 |-- edge_type: string (nullable = false)
```

```
In [25]: comments_questions_sdf.count()
```

```
Out[25]: 20268151
```

```
In [26]: comments_questions_sdf.show(5)
```

```
+-----+-----+-----+
|from_node|to_node|edge_type|
+-----+-----+-----+
|      4550|    4550|comment-on-question|
|       242|     184|comment-on-question|
|     4213|    4946|comment-on-question|
|        91|      91|comment-on-question|
|     2658|    1874|comment-on-question|
+-----+-----+-----+
only showing top 5 rows
```

```
In [27]: comments_questions_sdf.printSchema()
```

```
root
 |-- from_node: integer (nullable = true)
 |-- to_node: integer (nullable = true)
 |-- edge_type: string (nullable = false)
```

```
In [28]: graph_sdf.count()
```

```
Out[28]: 63497050
```

```
In [29]: graph_sdf.show(5)
```

```
+-----+-----+-----+
|from_node|to_node|edge_type|
+-----+-----+-----+
|        9|      8|  answer|
|         1|      1|  answer|
|       13|      1|  answer|
|       17|      1|  answer|
|       48|      2|  answer|
+-----+-----+-----+
only showing top 5 rows
```

```
In [30]: graph_sdf.printSchema()
```

```
root
 |-- from_node: integer (nullable = true)
 |-- to_node: integer (nullable = true)
 |-- edge_type: string (nullable = false)
```

1.5 Step 2.4

```
In [31]: # You may put any computations you need here
no_of_questions_sdf=answers_sdf.unionAll(comments_questions_sdf)
no_of_questions_sdf=no_of_questions_sdf.groupBy(no_of_questions_sdf['to_node'].alias('u
no_of_questions_sdf=no_of_questions_sdf.orderBy('count',ascending=False)
#no_of_questions_sdf.select(no_of_questions_sdf['to_node'].alias('user')).show(5)

In [32]: no_of_answers_sdf=answers_sdf[(answers_sdf.from_node)!=(answers_sdf.to_node)]
no_of_answers_sdf=no_of_answers_sdf.groupBy(no_of_answers_sdf['from_node'].alias('user'
no_of_answers_sdf=no_of_answers_sdf.orderBy('count',ascending=False)

In [33]: list_op=(comments_questions_sdf.select('to_node').subtract(answers_sdf.select('to_node'

In [34]: #sdf1=answers_sdf.groupBy(answers_sdf.from_node,answers_sdf.to_node).count()
#sdf2=sdf1
#sdf1.join(sdf2,(F.col(sdf1.to_node)==(F.col(sdf2.from_node))), (F.col(sdf1.to_node)==(F
answers_sdf.groupBy("from_node", "to_node").count()

#(answers_sdf.groupBy("from_node", "to_node").count().as("count1")).
# join(answers_sdf.groupBy("col1").agg(sum($"col3").as("sum_level1")), Seq("col1"))
```

```
Out[34]: DataFrame[from_node: int, to_node: int, count: bigint]
```

```
In [35]: answers_sdf.show(10)
```

```
+-----+-----+-----+
|from_node|to_node|edge_type|
+-----+-----+-----+
|      9|      8|  answer|
|      1|      1|  answer|
|     13|      1|  answer|
|     17|      1|  answer|
|     48|      2|  answer|
|     17|      1|  answer|
|     19|      9|  answer|
|     13|     23|  answer|
|     13|     11|  answer|
|     23|     23|  answer|
+-----+-----+-----+
```

only showing top 10 rows

1.6 Step 2.4.1 Results

```
In [36]: # TODO: output dataframe with top 10 users by number of questions
no_of_questions_sdf.show(10)
```

```
+-----+-----+
|  user|count|
+-----+-----+
|875317| 7995|
|  4653| 7652|
| 34537| 6991|
|117700| 6830|
| 39677| 6686|
|541686| 5618|
| 84201| 5528|
|859154| 5500|
| 65387| 5464|
|179736| 5439|
+-----+-----+
only showing top 10 rows
```

```
In [37]: # TODO: output top 10 users by number of answers to questions by distinct users
         no_of_answers_sdf.show(10)
```

```
+-----+-----+
|  user|count|
+-----+-----+
| 22656|32020|
|1144035|25146|
| 29407|20836|
| 548225|16939|
| 157882|16609|
| 115145|16503|
| 17034|15436|
| 100297|15017|
| 6309|14276|
| 34397|14013|
+-----+-----+
only showing top 10 rows
```

1.7 Step 2.4.2 Results

```
In [38]: # TODO: number of users whose questions have never been answered or commented on
         len(list_op)
```

```
Out[38]: 150681
```

1.8 Step 2.4.3 Results

```
In [39]: # TODO: top 10 pairs of users by mutual answers, along with the number of questions the
         anscount_df = answers_sdf.groupby('from_node' , 'to_node').count()
```



```

anscount_df1 = anscount_df.select(anscount_df['from_node'].alias("from_node_1"),
                                   anscount_df['to_node'].alias("to_node_1"),anscount_df)

anscount_df2 =  anscount_df.select(anscount_df['from_node'].alias("from_node_2"),
                                   anscount_df['to_node'].alias("to_node_2"),anscount_df)

concat_df = anscount_df1.join(anscount_df2 , ((anscount_df1['from_node_1'] ) == (anscount_df2['from_node_2'] ))
                                ((anscount_df1['to_node_1']) == (anscount_df2['from_node_2'])) ))
ans_count = concat_df['count_1'] + concat_df['count_2']
df = concat_df.withColumn("anscount" , ans_count)
df = df[df.from_node_1 != df.to_node_1]

df = df.select(df['from_node_1'].alias("user1"),
               df['to_node_1'].alias("user2"),df['anscount'].alias("anscount"))

df = df.sort('anscount' , ascending=False)
df.dropDuplicates()
df.show(10)

```

```

+-----+-----+-----+
| user1| user2|anscount|
+-----+-----+-----+
| 366797| 15168|      65|
| 15168| 366797|      65|
| 650492| 505088|      57|
| 505088| 650492|      57|
|1525840|1675891|      49|
|1675891|1525840|      49|
|2313718|2138752|      46|
|2138752|2313718|      46|
|1931641|1642617|      41|
|1642617|1931641|      41|
+-----+-----+-----+
only showing top 10 rows

```

2 Step 3

```

In [40]: # TODO: remove these, which just create dummy data
         #highest_indegree_sdf = spark.createDataFrame([{'node': 123, 'indegree': 4}])
         #highest_outdegree_sdf = spark.createDataFrame([{'node': 123, 'outdegree': 5}])

         # TODO: Fill in according to HW spec
         out_df = answers_sdf.groupby('from_node').count()
         out_df = out_df.select(out_df['from_node'].alias("node"),

```

```

                                out_df['count'].alias("outdegree"))
out_df = out_df.sort('count' , ascending=False)
highest_outdegree_sdf=out_df

in_df  = answers_sdf.groupby('to_node').count()
in_df = in_df.select(in_df['to_node'].alias("node"),
                    in_df['count'].alias("indegree"))

in_df = in_df.sort('count' , ascending=False)
highest_indegree_sdf=in_df

```

2.1 Step 3 Results

In [41]: `highest_indegree_sdf.show(5)`

```

+-----+-----+
| node|indegree|
+-----+-----+
|  4653|    5453|
| 39677|    4971|
| 34537|    4391|
|179736|    3716|
| 84201|    3710|
+-----+-----+
only showing top 5 rows

```

In [42]: `highest_outdegree_sdf.show(5)`

```

+-----+-----+
| node|outdegree|
+-----+-----+
| 22656|    32030|
|1144035|    25146|
| 29407|    20842|
| 548225|    16944|
| 157882|    16615|
+-----+-----+
only showing top 5 rows

```

2.2 Step 4

In [1]: *# TODO: insert code as you like*

```

In [66]: #
         # Step 4.1 Pre-processing

```

```

#
def spark_bfs(G, origins, max_depth):
    ##Your logic goes here
    schema = StructType([StructField("node", IntegerType(), True)])
    my_sdf = spark.createDataFrame(origins, schema)
    frontier = my_sdf
    visited = spark.createDataFrame([], frontier.schema)
    for i in range(max_depth+1):
        if i == 0:
            return_sdf = frontier
            return_sdf = return_sdf.withColumn('new_col', F.lit(i)).alias('depth')
        else:
            d1 = frontier.alias('f').join(G.alias('g'), F.col('f.node') == F.col('g.from_node'))
            d2 = d1.select(d1['to_node'].alias("node"))
            visited = visited.unionAll(frontier)
            frontier = d2
            frontier = frontier.join(visited, frontier.node == visited.node, 'leftanti')
            G = G.join(visited, G.to_node == visited.node, 'leftanti')
            temp_df = frontier
            temp_df = temp_df.withColumn('new_col', F.lit(i)).alias('depth')
            return_sdf = return_sdf.unionAll(temp_df)

    return return_sdf

```

2.3 Step 4.1

```

In [67]: # TODO: comment out this line once your code is ready
         #bfs_sdf = spark.createDataFrame([{'node': 123, 'depth': 1}, {'node': 456, 'depth': 2}])
         # TODO: enable this once your code is ready
         origin_map = [{'node': 4550}, {'node': 242}]
         bfs_sdf = spark_bfs(comments_questions_sdf, origin_map, 2)
         bfs_sdf.cache()
         bfs_sdf.count()

```

Out[67]: 397

```

In [2]: # TODO: insert code as you like

```

2.4 Step 4.1 Results

```

In [68]: bfs_sdf.show(10)

```

```

+-----+-----+
|  node|new_col|
+-----+-----+
|  4550|      0|
|   242|      0|
|1619254|     1|
|2332659|     1|

```

```
|5504881|      1|
|1139389|      1|
|1940564|      1|
| 818089|      1|
|3047450|      1|
|4773326|      1|
```

```
+-----+-----+
```

only showing top 10 rows

2.5 Step 4.2

```
In [71]: #
         # Step 4.2 Pre-processing
         def create_filtered_bfs_sdf(input_sdf):
             filtered_bfs_sdf = input_sdf[input_sdf.new_col == 2]
             filtered_bfs_sdf = filtered_bfs_sdf.groupBy('node').count()
             filtered_bfs_sdf = filtered_bfs_sdf[filtered_bfs_sdf['count'] > 1]
             filtered_bfs_sdf.cache()
             return filtered_bfs_sdf

         def friend_rec(input_sdf, graph_sdf):
             G1=filtered_bfs_sdf.select(filtered_bfs_sdf['node'].alias("from_node"))
             G2=filtered_bfs_sdf.select(filtered_bfs_sdf['node'].alias("to_node"))
             G3=G1.join(G2, (G1['from_node'] < G2['to_node']))
             temp_sdf=graph_sdf.select(graph_sdf['from_node'],graph_sdf['to_node'])
             G3=G3.subtract(temp_sdf)
             temp_sdf=graph_sdf.select(graph_sdf['to_node'].alias('from_node'),graph_sdf['from_n
             friend_recommendations_sdf=G3.subtract(temp_sdf)
             return friend_recommendations_sdf

In [72]: # TODO: insert code as you like

In [73]: # TODO: comment this line out when your function works
         #friend_recommendations_sdf = spark.createDataFrame([\
                                                     # {'from_node': 123, 'to_node': 456}
                                                     # {'from_node': 456, 'to_node': 123}]

         # TODO: enable this when your function works
         filtered_bfs_sdf = create_filtered_bfs_sdf(bfs_sdf)
         friend_recommendations_sdf = friend_rec(filtered_bfs_sdf, comments_questions_sdf)
         friend_recommendations_sdf.cache()
         friend_recommendations_sdf.count()

Out[73]: 902

In [74]: friend_recommendations_sdf.show(5)

+-----+-----+
|from_node|to_node|
```

```

+-----+-----+
|    21918| 453447|
|    21918|4204628|
|    46646| 267679|
|    59017|1187554|
|    503032|2269511|
+-----+-----+
only showing top 5 rows

```

2.6 Step 4.2 Results

```
In [75]: friend_recommendations_sdf.show()
```

```

+-----+-----+
|from_node|to_node|
+-----+-----+
|    21918| 453447|
|    21918|4204628|
|    46646| 267679|
|    59017|1187554|
|    503032|2269511|
|    46646|1127460|
|    55503| 104015|
|    60602|1187554|
|    267679|1061543|
| 1218595|4800193|
|    27483|  46646|
|   104015|4595831|
|   211452|1894566|
| 1038179|1187554|
|    60602|  63775|
|   282194| 503032|
|   830423|4899760|
|    21918|  59017|
|    21918|1187554|
|    59017| 830423|
+-----+-----+
only showing top 20 rows

```

2.7 Step 4.3: Graph visualization

2.7.1 Once you have executed the cells in Step 4.2 and you have `friend_recommendations_sdf`, let's create `friend_recommendations_df` using `toPandas()`. This creates an in-memory DataFrame that we can use to build the graph. Here we have used ('from_node','to_node') as column names in `friend_recommendations_sdf`, please change it to what you have used in yours.

```
In [76]: import networkx as nx
# TODO: create friend_graph NetworkX graph from friend_recommendations_df from friend_r
n_sdf=friend_recommendations_sdf.toPandas()
friend_graph=nx.from_pandas_dataframe(n_sdf, 'from_node', 'to_node')
```

2.8 Step 4.3 Results

```
In [77]: print ("Number of nodes (characters) in this graph is", friend_graph.order()) # number
print ("Number of edges in this graph is", len(friend_graph.edges())) # number of edges
print ("Graph diameter is", nx.diameter(friend_graph)) # maximum eccentricity
```

Number of nodes (characters) in this graph is 43

Number of edges in this graph is 902

Graph diameter is 2

```
In [78]: %matplotlib inline
nx.draw(friend_graph)
```

